

The background is a dark blue gradient with a subtle pattern of small white dots. Overlaid on the left side are several concentric circular patterns. A prominent circular scale with tick marks and numbers (140, 150, 160, 170, 180, 190, 200, 210, 220, 230, 240, 250, 260) is visible. Other circular elements include dashed lines, solid lines, and arrows indicating clockwise or counter-clockwise rotation.

MEMORY MANAGEMENT IN PROGRAMMING LANGUAGES

LECTURE # 5

WHAT IS MEMORY MANAGEMENT?

- Memory management is the process of controlling and coordinating the way a software application access **computer memory**
- When a software runs on a target Operating system on a computer it needs access to the computers **RAM**(Random-access memory) to:
 - load its own **bytecode** that needs to be executed
 - store the **data values** and **data structures** used by the program that is executed
 - load any **run-time systems** that are required for the program to execute

REVIEW DEFINITIONS

- **Method**: any subprogram (function, procedure, subroutine) – depends on language terminology.
- **Environment** of an **active method**: the variables it can currently access plus their addresses (a set of ordered pairs)
- **State** of an active method: variable/value pairs

THREE CATEGORIES OF MEMORY (FOR DATA STORE)

- **Static**: storage requirements are known prior to run time; lifetime is the entire program execution
- **Run-time stack**: memory associated with active functions
 - Structured as stack frames (activation records)
- **Heap**: dynamically allocated storage; the least organized and most dynamic storage area

STATIC DATA MEMORY

- Simplest type of memory to manage.
- Consists of anything that can be completely determined at compile time; *e.g.*, global variables, constants (perhaps), code.
- Characteristics:
 - Storage requirements known prior to execution
 - Size of static storage area is constant throughout execution

RUN-TIME STACK

- The stack is a contiguous memory region that grows and shrinks as a program runs.
- Its purpose: to support method calls
- It grows (storage is allocated) when the **activation record** (or **stack frame**) is pushed on the stack at the time a method is called (activated).
- It shrinks when the method terminates and storage is de-allocated.

RUN-TIME STACK

- The stack frame has storage for local variables, parameters, and return linkage.
- The size and structure of a stack frame is known at compile time, but actual contents and time of allocation is unknown until runtime.
- How is variable lifetime affected by stack management techniques?

HEAP MEMORY

- Heap objects are allocated/deallocated dynamically as the program runs (not associated with specific event such as function entry/exit).
- The kind of data found on the heap depends on the language
 - Strings, dynamic arrays, objects, and linked structures are typically located here.
 - Java and C/C++ have different policies.

HEAP MEMORY

- Special operations (e.g., **malloc**, **new**) may be needed to allocate heap storage.
- When a program deallocates storage (**free**, **delete**) the space is returned to the heap to be re-used.
- Space is allocated in variable sized blocks, so deallocation may leave “holes” in the heap (fragmentation).
 - Compare to deallocation of stack storage

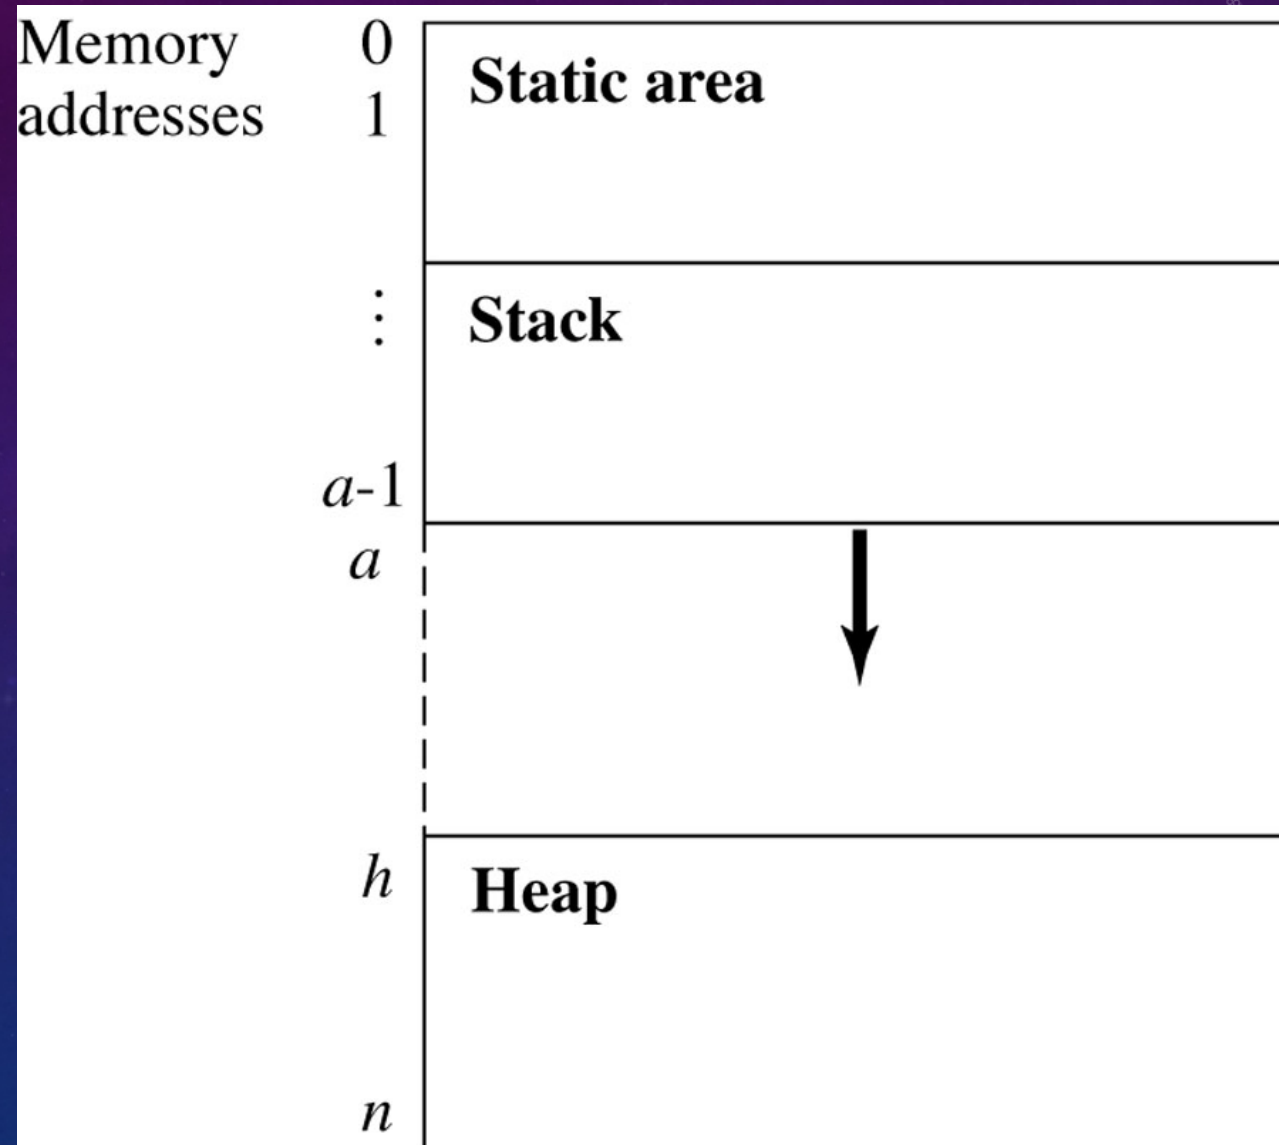
HEAP MANAGEMENT

- Some languages (e.g. C, C++) leave heap storage deallocation to the programmer
 - `delete`
- Others (e.g., Java, Perl, Python, list-processing languages) employ **garbage collection** to reclaim unused heap space.

The Structure of Run-Time Memory

Figure 11.1

These two areas grow towards each other as program events require.



STACK OVERFLOW

- The following relation must hold:
 $0 \leq a \leq h \leq n$
- In other words, if the stack top bumps into the heap, or if the beginning of the heap is greater than the end, there are problems!

ARRAY OUT-OF-BOUNDS VIOLATIONS IN C

- No run-time type-checking is taking place as a C program executes.
- No-one is tracking array bounds violations

HEAP STORAGE STATES

- For simplicity, we assume that memory words in the heap have one of three states:
 - Unused: not allocated to the program yet
 - Undef: allocated, but not yet assigned a value by the program
 - Contains some actual value

HEAP MANAGEMENT FUNCTIONS

- *new* returns the start address of a block of k words of unused heap storage and changes the state of the words from *unused* to *undef*.
 - $n \leq k$, where n is the number of words of storage needed; e.g., suppose a Java class Point has data members x,y,z which are floats.
 - If floats require 4 bytes of storage, then
Point firstCoord = new Point()
calls for 3 X 4 bytes (at least) to be allocated and initialized to some predetermined state.

HEAP OVERFLOW

- **Heap overflow** occurs when a call to *new* occurs and the heap does not have a contiguous block of k unused words
- So *new* either fails, in the case of heap overflow, or returns a pointer to the new block

HEAP MANAGEMENT FUNCTIONS

- *delete* returns a block of storage to the heap
- The status of the returned words are returned to *unused*, and are available to be allocated in response to a future *new* call.
- One cause of heap overflow is a failure on the part of the program to return unused storage.

HEAP ALLOCATION

- Heap space isn't necessarily allocated and deallocated from one end (like the stack) because the memory is not allocated and deallocated in a predictable (first-in, first-out or last-in, first-out) order.
- As a result, the location of the specific memory cells depends on what is available at the time of the request.

MEMORY ATTRIBUTES

- Memory to store data in programming languages has the following lifecycle
 - Allocation: When the memory is allocated to the program
 - Lifetime: How long allocated memory is used by the program
 - Recovery: When the system recovers the memory for reuse

MEMORY CLASSES

- **Static memory – Usually at a fixed address**
 - Lifetime – The execution of program
 - Allocation – For entire execution
 - Recovery – By system when program terminates
 - Allocator – Compiler
- **Automatic (LIFO) memory – Usually on a stack**
 - Lifetime – Activation of method using that data
 - Allocation – When method is invoked
 - Recovery – When method terminates
 - Allocator – Typically compiler, sometimes programmer

MEMORY CLASSES ...CONT

- Dynamic memory – Addresses allocated on demand in an area called the heap
 - Lifetime – As long as memory is needed
 - Allocation – Explicitly by programmer, or implicitly by compiler
 - Recovery – Either by programmer or automatically (when possible and depends upon language)
 - Allocator – Manages free/available space in heap

MEMORY MANAGEMENT IN C

- Local variables live on the stack
 - Allocated at function invocation time
 - Deallocated when function returns
 - Storage space reused after function returns
- Space on the heap allocated with malloc()
 - Must be explicitly freed with free()
 - Called explicit or manual memory management
 - Deletions must be done by the user

MEMORY MANAGEMENT

- Computer programs need to allocate memory to store data values and data structures.
- Memory is also used to store the program itself and the run-time system needed to support it.
- If a program allocates memory and never frees it, and that program runs for a sufficiently long time, eventually it will run out of memory.
- Even in the presence of virtual memory, memory consumption is still a major issue because it is considerably less efficient to access virtual memory than to access physical memory

MANUAL & AUTOMATIC MEMORY MANAGEMENT

- **Automatic memory management**
- **Which ask the programmer to allocate and free memory manually.**
 - The C language requires the programmer to implement memory management each time, for each application program.
 - Modern programming languages such as Java, C#, Caml, Cyclone and Ruby provide automatic memory management with garbage collection

MANUAL MEMORY MANAGEMENT

- In C, where there is no garbage collector, the programmer must allocate and free memory explicitly.
 - The key functions are malloc and free.
- The malloc function takes as a parameter the size in bytes of the memory area to be allocated.
- The size of a type can be obtained using sizeof.
- The resulting area of memory does not represent a value of the correct type, so it then needs to be cast to the correct type.

```
p = (Type_t*) malloc(sizeof(Type_t));
```


DANGLING POINTER PROBLEM

- A significant problem with manual memory management is that it is possible to attempt to use a pointer after it has been freed.
- This is known as the dangling pointer problem. Dangling pointer errors can arise whenever there is an error in the control flow logic of a program.
- This can lead to allocation, use and deallocation happening in the wrong order in some circumstances
- Use before allocation may be a fatal run-time error. Use after deallocation is not always fatal. Neither of these is a good thing.

DANGLING POINTER

```
{ int *x = ...malloc();  
  free(x);  
  *x = 5; /* oops! */  
}
```


SPACE/MEMORY LEAK

- Another potential problem of manual memory management is not remembering to free allocated memory when it should be freed.
- The reference to an allocated area of memory can be lost when a variable in a block-structured language goes out of scope.
- This problem is perhaps more subtle than the dangling pointer problem because it may only become manifest for long-running applications.
- When memory is lost and cannot be reclaimed we term this a space leak. Space cannot be lost forever without reaching the limit on the available memory.
- sA long-running program with a space leak will eventually crash.

MEMORY LEAK

```
{ int *x = (int *) malloc(sizeof(int)); }
```



MAY FREE SOMETHING TWICE

```
{ int *x = ...malloc(); free(x); free(x); }
```

FRAGMENTATION

- Another memory management problem
- Example sequence of calls

```
allocate(a);  
allocate(x);  
allocate(y);  
free(a);  
allocate(z);  
free(y);  
allocate(b);
```



⇒ Not enough contiguous space for b

AUTOMATIC MEMORY MANAGEMENT

- Primary goal: automatically reclaim dynamic memory.
- Secondary goal: also avoid fragmentation

